

A Sequential Hybrid Forecasting System for Gold Returns: Integrating ARIMAX-GARCH with Machine Learning for Adaptive Tail-Risk Management

Pham Thanh Long

Student ID: 11230562

Class: Data Science in Economics and Business 65B

Course: Time Series Analysis

Department of Mathematical Economics

National Economics University, Hanoi, Vietnam

Email: 11230562@st.neu.edu.vn

May 2026

Abstract

The global financial landscape between 2019 and 2026 has been defined by unprecedented macroeconomic volatility, characterized by escalating geopolitical conflicts, supply chain reconfigurations, and aggressive monetary policy pivots by major central banks. In this environment of heightened systemic uncertainty, gold has re-emerged as a critical safe-haven asset and an essential vehicle for inflation hedging and capital preservation. However, modeling the trajectory of gold returns under these high-stress regimes poses severe challenges to traditional quantitative frameworks, as asset dynamics increasingly exhibit heavy-tailed distributions, non-linear shocks, and conditional heteroskedasticity. To address these structural anomalies, I develop an empirical research design that evaluates the integration of macro-econometric modeling with a comparative cohort of machine learning algorithms, utilizing a daily transaction dataset spanning from January 2019 through April 2026.

Instead of relying on a single unconstrained model, I implement a multi-stage sequential framework where an econometric ARIMAX specification first filters foundational linear macro-trends, binding gold returns to a filtered vector of global leading indicators: the U.S. Dollar Index (DXY), market anxiety metrics (VIX), real capital costs (US10Y), energy inflation (Oil), and equity market benchmarks (SPX). To capture the remaining conditional heteroskedasticity, the linear residual stream is modeled through an autoregressive GARCH(1,1) specification under heavy-tailed conditional distribution constraints. Crucially, to identify the optimal mechanism for non-linear error correction, I execute a rigorous comparative analysis across three alternative machine learning architectures: eXtreme Gradient Boosting (XGBoost), Random Forest (RF), and Support Vector Regression (SVR). Each algorithm incorporates autoregressive error lags and technical indicators, utilizing the explicit GARCH conditional standard deviation ($\sigma_{t,\text{GARCH}}$) to guide its respective statistical mappings.

The empirical verification reveals distinct trade-offs in predictive performance and structural stability among the three machine learning candidates under review. Beyond statistical forecasting accuracy, I demonstrate the practical utility of the system through an adaptive risk management backtesting simulation. By embedding the time-varying GARCH variance parameters derived from the optimal hybrid model directly into the portfolio layout, I construct dynamic Value-at-Risk (VaR) bands that automatically expand during periods of market stress and contract during low-risk regimes. This methodology delivers a comprehensive, single-author quantitative framework that successfully balances directional machine learning execution with proactive capital preservation and automated stop-loss management.

Keywords— Hybrid System, ARIMAX-GARCH, XGBoost, Random Forest, Support Vector Regression, Volatility Clustering, Value-at-Risk, Gold Forecasting.

1 Introduction

1.1 Macroeconomic Instability and the Asset Profile of Gold

The global financial ecosystem from January 2019 through April 2026 has been characterized by intense macroeconomic instability and structural transitions. This volatile multi-year timeline incorporates consecutive systemic disruptions, including a global pandemic, escalating geopolitical conflicts across Europe and the Middle East, widespread supply chain reconfigurations, and an aggressive monetary policy pivot by major central banks shifting from long-term loose liquidity expansion to unprecedented quantitative tightening.

Within this historical landscape of elevated stress, gold has re-asserted its foundational institutional function as an ultimate safe-haven asset, a sovereign store of value, and an organic vehicle for capital preservation. As cross-asset correlation matrices breakdown and fiat purchasing power undergoes rapid depreciation, global capital consistently migrates toward gold, positioning its price trajectory as a critical real-time indicator of systemic anxiety and financial risk transmission.

1.2 Methodological Complications in Financial Time-Series

Despite the clear importance of gold price movements for global asset allocation, executing reliable short-to-medium-term forecasting poses significant challenges for quantitative finance. Daily gold log-return series do not evolve along clean, static linear pathways. Instead, under modern stress regimes, they are characterized by distinct mathematical anomalies: heavy-tailed conditional distributions, sharp asymmetrical shocks, and intense time-varying variance propagation—commonly defined as volatility clustering.

As developed extensively in the quantitative literature, traditional forecasting frameworks generally fail to resolve these properties simultaneously due to an operational trade-off. Linear econometric models (such as ARIMAX variations) preserve solid economic interpretability by linking returns directly to fundamental macroeconomic drivers like interest rates or currency shifts, yet they remain structurally blind to non-linear momentum and time-varying variance. Conversely, pure non-linear machine learning algorithms excel at mapping complex high-frequency updates, but they function as unconstrained "black boxes" that systematically ignore underlying economic principles and remain highly vulnerable to out-of-sample over-fitting on raw market noise.

1.3 Bridging the Literature Gaps: A Multi-Stage Strategy

As identified in the comprehensive review of previous studies, traditional quantitative literature suffers from three critical boundaries during periods of severe financial distress:

- **The Volatility Blind-Spot:** Most hybrid frameworks pass econometric residuals directly into machine learning sequences without addressing conditional heteroskedasticity, leaving the algorithm blind to sudden shifts in the underlying market risk regime.
- **The Absence of Competing Algorithmic Cohorts:** Existing structures typically rely on a single arbitrarily chosen model without executing a systematic, simultaneous comparison across fundamentally different machine learning classes within an identical macro environment.
- **The Isolation of Risk Management:** Quantitative papers focus almost exclusively on minimizing abstract statistical error metrics, failing to connect directional machine learning output with actionable portfolio insulation mechanisms.

To resolve these distinct empirical gaps, my independent study designs an integrated, sequential three-tier hybrid pipeline. The structural framework enforces a strict division of labor where an econometric ARIMAX base first filters long-term linear macro-trends, an autoregressive GARCH(1,1) model isolates conditional volatility clustering, and a competitive cohort of alternative machine learning regressors executes precise non-linear error correction.

1.4 Research Question and Core Objectives

The central objective of my independent study is to answer a fundamental quantitative research question: *Can an integrated hybrid architecture—which explicitly injects econometric macro-dependencies and GARCH volatility parameters into a competitive machine learning framework—outperform standalone forecasting benchmarks while successfully constructing an adaptive risk-insulation mechanism during prolonged crisis cycles?*

To systematically answer this question, I establish three operational objectives:

1. To build a macro-econometric ARIMAX baseline that binds gold log-returns to an empirical vector of global leading indicators, including the U.S. Dollar Index (DXY), market anxiety metrics (VIX), capital costs (US10Y), energy metrics (Oil), and equity returns (SPX).
2. To fit an autoregressive GARCH(1,1) specification under heavy-tailed distribution constraints to isolate the explicit conditional standard deviation ($\sigma_{t,\text{GARCH}}$) of the linear error stream.
3. To execute a rigorous comparative experiment across three distinct algorithmic classes—eXtreme Gradient Boosting (**XGBoost**), Random Forest (**RF**), and Support Vector Regression (**SVR**)—by feeding the GARCH volatility series directly into their feature space to determine the optimal non-linear error-correction mechanism.

2 Literature Review

2.1 Macroeconomic Transmission Channels and Gold Dynamics

The foundational literature governing commodity pricing establishes that gold prices do not evolve within an economic vacuum; instead, they are driven by complex macroeconomic transmission vectors. Early asset pricing theories position gold as the ultimate sovereign store of value and an organic hedge against systemic fiat currency depreciation. Academic consensus, represented by standard financial models, confirms a strong negative correlation between the U.S. Dollar Index (DXY) and gold returns, driven by the currency denomination effect.

Furthermore, empirical frameworks focusing on cost-of-capital vectors emphasize that gold, as a non-yielding asset, exhibits high sensitivity to changes in real interest rates, traditionally proxied by long-term government bond yields such as the US10Y. During periods of aggressive central bank interventions and quantitative tightening cycles—as observed intensely within the 2019–2026 market regime—these macro-dependencies undergo severe structural shifts. Additionally, commodity economic literature paths a direct link between energy-led inflation (primarily driven by Crude Oil shocks) and gold return expansions, where gold acts as a physical shield against cost-push inflation cascades.

2.2 The Evolution of Linear Econometric and Volatility Modeling

From a methodological perspective, the classical paradigm for financial time-series forecasting has long been dominated by the linear Box-Jenkins framework. The Autoregressive Integrated Moving Average (ARIMA) model and its exogenous extension (ARIMAX) have been extensively utilized due to their strong short-term forecasting capacity, ease of implementation, and asymptotic stability when handling non-stationary data. ARIMAX models successfully preserve economic interpretability by directly incorporating fundamental macro-drivers into the conditional mean equation.

However, the primary limitation of pure ARIMAX architectures is their fundamental assumption of constant variance, which makes them incapable of capturing the non-linear, time-varying volatility dynamics inherent in financial return streams. To resolve this structural flaw, Engle (1982) introduced the Autoregressive Conditional Heteroskedasticity (ARCH) model [cite: 714], which Bollerslev (1986) subsequently generalized into the GARCH framework[cite: 716].

The integration of ARIMA-type mean equations with GARCH variance components allowed researchers to model conditional heteroskedasticity and volatility clustering effectively[cite: 7]. Standard univariate ARIMA-GARCH models became a benchmark in commodity forecasting, proving that modeling the residual variance under non-Gaussian distributions (such as the Student- t or Skewed Student- t) significantly improves tail-risk coverage and interval estimation accuracy[cite: 7, 12, 13].

2.3 Algorithmic Machine Learning in Residual Error Correction

With the advent of high-dimensional computing, the quantitative finance literature shifted toward non-linear Machine Learning (ML) architectures to capture complex patterns that econometric models fail to filter. Researchers successfully deployed tree-based ensemble methods and kernel-based optimizers to model asset trajectories. Among these algorithms, eXtreme Gradient Boosting (XGBoost), Random Forest (RF), and Support Vector Regression (SVR) emerged as superior forecasting tools due to their capacity to map high-frequency non-linear updates and handle structural data breaks without rigid parametric constraints.

Despite their predictive accuracy, purely algorithmic approaches face severe criticism in financial economics. Because they lack an innate economic structure, unconstrained machine learning models often over-fit on raw price noise, leading to unstable out-of-sample predictions when transitioning between different market regimes.

To bridge this division, a secondary branch of literature emerged focusing on hybrid architectures. These frameworks operate on a sequential "two-phase" or "multi-stage" process: a linear econometric base model filters structural trends first, isolating a residual error stream that contains purely non-linear signals. This residual series is then passed to a machine learning regressor for non-linear error correction. This combination balances the structural interpretability of econometrics with the flexible precision of machine learning algorithms.

2.4 Identified Research Gaps and My Structural Positioning

Although the existing literature provides strong evidence supporting hybrid time-series modeling, several critical gaps remain unresolved, particularly when exposed to highly unstable financial environments such as the January 2019–April 2026 market cycle:

1. **The Volatility Blind-Spot in Machine Learning Feature Spaces:** Most sequential hybrid models pass the linear residual series (E_t) directly into the machine learning algorithm, ignoring the time-varying variance properties of the error stream. By leaving out conditional volatility (σ_t) from the machine learning feature space, the algorithm remains blind to changes in the underlying risk regime, often leading to poor performance during sudden market panics.
2. **The Lack of Rigorous Comparative Cohorts:** Existing hybrid gold-forecasting literature typically relies on a single chosen machine learning algorithm (e.g., a simple neural network or a single tree model) without executing a systematic, simultaneous experiment across competing algorithmic classes (such as boosting vs. bagging vs. kernel optimization) within an identical macroeconomic framework.
3. **The Separation of Forecasting and Actionable Risk Management:** Traditional hybrid papers focus almost entirely on reducing statistical error metrics (such as minimizing RMSE or MAPE), failing to connect the directional machine learning forecast with proactive, adaptive portfolio insulation mechanisms like dynamic Value-at-Risk (VaR) quantiles.

My independent study is designed specifically to fill these distinct gaps. Instead of running a basic univariate or unconstrained machine learning model, I implement a structural three-tier sequential pipeline.

First, my model establishes macroeconomic grounding by utilizing an ARIMAX base to filter out linear macro trends linked directly to global factors (DXY, VIX, US10Y, Oil, SPX). Second, I address

the volatility blind-spot by fitting a GARCH(1,1) model to the linear errors, explicitly routing the extracted time-varying conditional standard deviation ($\sigma_{t,\text{GARCH}}$) into the machine learning feature matrix. This allows me to conduct a rigorous comparative experiment across three alternative machine learning algorithms (**XGBoost**, **Random Forest**, **SVR**) to isolate the optimal error-correction mechanism. Finally, I translate these integrated statistical outputs into an actionable risk-adjusted risk management asset allocation framework by constructing adaptive VaR bands that dynamically scale portfolio exposure to ensure capital protection during crisis regimes.

3 Data and Methodology

3.1 Data Integration and Operational Metrics

To evaluate the predictive capacity of the proposed hybrid architecture, I construct a daily financial dataset spanning from January 2019 through April 2026. The temporal boundaries of this sample capture multiple distinct structural regimes: the pre-pandemic stable environment, the high-liquidity expansionary phase of 2021, and the quantitative tightening cycles alongside intense geopolitical shocks marking the 2022–2026 era.

The primary target space is defined by the price of Gold Futures. To map the multi-dimensional macroeconomic transmission channels that govern gold’s safe-haven and inflation-hedging profile, I integrate five daily exogenous indicators. To eliminate unit-root non-stationarity, prevent spurious regressions, and stabilize conditional variance, all raw price levels are transformed into continuously compounded daily log-returns:

$$r_t = 100 \times \ln \left(\frac{P_t}{P_{t-1}} \right) \quad (1)$$

Where P_t represents the asset closing price on day t .

Conversely, applying rate-of-growth transformations to indicators that are natively expressed as percentages or bounded point distributions creates structural distortions and arithmetic boundary anomalies. Therefore, macro-financial variables that are natively expressed as interest rates, yields, or index points are operationalized using the absolute first-difference operator (Δ). Mathematically, the first-difference filter isolates the absolute period-to-period change and is formulated as follows:

$$\Delta X_t = X_t - X_{t-1} \quad (2)$$

Where X_t denotes the raw nominal value of the indicator at session t . For the 10-Year U.S. Treasury Bond Yield, this specific filter extracts the marginal day-to-day shift directly in basis points (Δ). This represents the pure change in the opportunity cost of holding non-yielding safe-haven assets like gold. For the CBOE Volatility Index (VIX), the absolute first difference measures the immediate velocity of shifts in market panic without generating mathematical errors near zero boundaries.

The complete systemic variable layout, along with their exact operational configurations and pythonic identifiers implemented within the dynamic backtesting codebase, is structured comprehensively in Table 1.

Table 1: Systemic Variable Specifications and Operational Transformations

Variable Class	Economic Indicator	Transformation Type	Unit
Target (Y_t)	Gold Price	Daily Log-Return	Percentage (%)
Exogenous ($X_{1,t}$)	U.S. Dollar Index	Daily Log-Return	Percentage (%)
Exogenous ($X_{4,t}$)	WTI Crude Oil Spot	Daily Log-Return	Percentage (%)
Exogenous ($X_{5,t}$)	S&P 500 Index	Daily Log-Return	Percentage (%)
Exogenous ($X_{2,t}$)	CBOE Volatility Index	Absolute First-Difference	Index Points (Δ)
Exogenous ($X_{3,t}$)	10-Year Treasury Yield	Absolute First-Difference	Index Points (Δ)




Figure 1: Time-series trajectories of transformed stationary target and exogenous variables.

3.2 Pre-Estimation Econometric Diagnostic Architecture

To safeguard the downstream dynamic empirical infrastructure against structural invalidity and model misspecification, the transformed dataset is subjected to a comprehensive sequence of econometric diagnostic filters. Classical quantitative modeling frequently overlooks these fundamental statistical attributes, resulting in overfitted mathematical weights that collapse when tested out-of-sample. This architecture establishes the asymptotic validity of the target return and exogenous predictors through four formal testing frameworks.

3.2.1 Stationarity and Unit Root Inference (ADF Test)

To prevent the mathematical catastrophe of "Spurious Regression"—wherein highly correlated deterministic trends yield artificially inflated R^2 metrics and invalid t -statistics—establishing structural stationarity is a mandatory prerequisite. A stochastic process is defined as stationary if its joint probability distribution is time-invariant, implying a constant mean, finite variance, and an autocovariance structure dependent strictly on temporal displacement lags rather than absolute time boundaries. To

empirically confirm this property, the Augmented Dickey-Fuller (ADF) unit-root test is executed using the following parametric autoregressive specification:

$$\Delta y_t = \alpha + \beta t + \gamma y_{t-1} + \sum_{i=1}^p \delta_i \Delta y_{t-i} + \varepsilon_t \quad (3)$$

where α is the structural drift intercept, t represents a continuous linear time trend, Δ is the first-difference operator, and ε_t is an independent and identically distributed (i.i.d.) white noise innovation vector. The statistical hypothesis tests the presence of a unit root via a one-sided boundary:

$$\begin{aligned} H_0 : \gamma &= 0 && \text{(Presence of a unit root, implying non-stationarity)} \\ H_1 : \gamma &< 0 && \text{(Mean-reverting stochastic trajectory, implying stationarity)} \end{aligned}$$

The optimal lag structure p is determined dynamically for each variable by minimizing the Akaike Information Criterion (AIC) to exhaust residual serial dependence.

3.2.2 Normality and Tail Diagnostics (Jarque-Bera Test)

Standard ordinary least squares estimations rely heavily on the assumption that structural residuals conform to a normal Gaussian distribution. To dissect the precise geometrical topology of the COMEX Gold return distribution (Y_t) and its macro drivers, the parametric Jarque-Bera (JB) test is deployed. The JB statistic measures the joint deviation of a sample's skewness and excess kurtosis from theoretical Gaussian boundaries, formulated as:

$$JB = \frac{n}{6} \left(S^2 + \frac{(K-3)^2}{4} \right) \quad (4)$$

where n represents the operational sample size, S denotes sample skewness ($S = E[(y_t - \mu)^3]/\sigma^3$), and K represents sample kurtosis ($K = E[(y_t - \mu)^4]/\sigma^4$). In a theoretical normal distribution, $S = 0$ and $K = 3$ (Excess Kurtosis = 0), collapsing the test statistic to zero.

3.2.3 Serial Autocorrelation Diagnostics (Ljung-Box Test)

To evaluate whether past return observations retain a temporal memory that could indicate market inefficiencies or introduce structural biases, I investigate the joint serial correlation profile using the Ljung-Box Q -test. This framework determines whether a group of autocorrelations up to a designated lag length m differs significantly from a zero vector, structured as follows:

$$Q_m = n(n+2) \sum_{k=1}^m \frac{\rho_k^2}{n-k} \quad (5)$$

where n is the total number of operational trading sessions, and ρ_k represents the sample autocorrelation coefficient at lag k . The null hypothesis (H_0) states that the series behaves as a white noise process up to lag m . In this diagnostic pipeline, the temporal dependence is monitored at a short-term milestone of lag $m = 5$.

3.2.4 Conditional Heteroskedasticity Inference (ARCH-LM Test)

A standard linear regression framework assumes homoskedastic error variance ($E[\varepsilon_t^2] = \sigma^2$). However, financial time series frequently exhibit time-varying volatility regimes. To test for the presence of Autoregressive Conditional Heteroskedasticity (ARCH), I estimate an auxiliary ARIMAX(1,0,1) model mapping gold returns against the exogenous macroeconomic matrix ($\mathbf{X}_t$). The extracted residuals (ε_t) are then evaluated using the Lagrange Multiplier (LM) framework based on the following squared residual regression:

$$\varepsilon_t^2 = \gamma_0 + \gamma_1 \varepsilon_{t-1}^2 + \gamma_2 \varepsilon_{t-2}^2 + \cdots + \gamma_q \varepsilon_{t-q}^2 + \nu_t \quad (6)$$

The statistical test targets the joint significance of lagged squared residuals up to order $q = 5$:

$$H_0 : \gamma_1 = \gamma_2 = \dots = \gamma_q = 0 \quad (\text{Homoskedastic residuals, no ARCH effects})$$

$$H_1 : \text{At least one } \gamma_i \neq 0 \quad (\text{Presence of conditional heteroskedasticity})$$

3.2.5 Empirical Diagnostic Findings and Discussion

The comprehensive pre-estimation diagnostic metrics computed across all system variables are systematically structuralized and presented in Table 2.

Table 2: Empirical Pre-Estimation Econometric Diagnostic Matrix

Variable	ADF Stat	ADF P -value	Skewness	Excess Kurtosis	LB P (Lag 5)	ARCH-LM P
Gold_Return	-18.0388	0.0000	-0.9725	10.6019	0.4118	0.0000
DXY_Return	-19.2846	0.0000	-0.2040	1.8070	0.7781	0.0000
VIX_Diff	-14.7910	0.0000	1.9076	29.4194	0.0000	0.0000
Oil_Return	-7.3738	0.0000	0.0228	24.0765	0.0002	0.0000
SPX_Return	-13.0303	0.0000	-0.6585	15.6392	0.0000	0.0000
US10Y_Diff	-32.0466	0.0000	-0.0664	1.7936	0.1133	0.0000

The empirical findings yield critical econometric insights that directly dictate the sequential modeling framework of the proposed hybrid architecture. First, the ADF test statistics are highly negative across all indicators, spanning from -7.3738 to -32.0466 , with all corresponding P -values converging asymptotically to 0.0000. This robust rejection of the null hypothesis at the strict 1% significance level provides empirical proof that all operationalized series are stationary processes integrated of order zero, or $I(0)$. This uniform stationarity eliminates the mathematical risk of "Spurious Regression" and justifies modeling the dynamic mean equations using an ARIMAX formulation without structural differencing.

Second, the structural morphology parameters exhibit severe deviations from theoretical Gaussian properties. The continuous daily log-returns of Gold possess a prominent negative skewness of -0.9725 and a substantial excess kurtosis of 10.6019, confirming a heavy-tailed (*leptokurtic*) trajectory. More dramatically, the CBOE Volatility Index shifts (VIX_Diff) and Crude Oil returns (Oil_Return) reveal extreme excess kurtosis values of 29.4194 and 24.0765, respectively. Paired with the complete rejection of normality by the Jarque-Bera test (P -value = 0.0000), the system exhibits an empirical distribution dominated by extreme outliers and structural asymmetric shocks. This definitive evidence exposes the limitations of classical linear models, which undercalculate tail-risk thresholds, and establishes the necessity of employing regularized non-linear machine learning ensembles (XGBoost, Random Forest, SVR) to absorb uncaptured structural residuals.

Third, the Ljung-Box Q -test at lag 5 uncovers a heterogeneous temporal memory structure. The target asset, Gold (P -value = 0.4118), and the opportunity cost driver, US10Y (P -value = 0.1133), exhibit no significant linear serial autocorrelation at short-term horizons, indicating that past states of these specific series do not provide strong linear signals. Conversely, VIX, Crude Oil, and the S&P 500 display intense temporal dependency (P -value < 0.01). This split configuration implies that a static time-series mean equation is insufficient, validating the integration of comprehensive macroeconomic transmission forces to harvest directional predictability.

Finally, the Autoregressive Conditional Heteroskedasticity Lagrange Multiplier (ARCH-LM) test establishes a uniform diagnostic profile across the entire variable space. Every indicator logs an ARCH-LM P -value of 0.0000, soundly rejecting the homoskedasticity null hypothesis. This diagnostic confirms the presence of "Volatility Clustering"—where high-volatility regimes cluster together in time—rendering traditional Ordinary Least Squares (OLS) estimations inefficient. This mathematical reality validates the inclusion of a dynamic GARCH(1,1) component to filter the conditional standard deviation (σ_t) before proceeding to feature matrix expansion.

3.2.6 Multivariate Distribution Analysis




Figure 2: Empirical probability distributions (KDE and Histograms) versus theoretical Gaussian normal benchmarks across the systemic macro-financial variable layout.

To complement the formal econometric testing matrix, a visual multi-plot dashboard matrix is constructed in Figure 2. This graphic infrastructure maps the empirical probability density functions—operationalized through non-parametric Kernel Density Estimations (KDE) and underlying frequency histograms—directly against theoretical Gaussian normal distributions parameterized specifically by each localized sample’s empirical mean (μ) and standard deviation (σ).

The visual topology of the system variables provides immediate, unambiguous validation of the parametric violations recorded in Table 2. Across all six panels, the empirical distributions (depicted by the solid blue lines and underlying shaded bars) display a profound deviation from the theoretical Gaussian bell curve (represented by the dashed crimson line). This structural divergence manifests primarily through two distinct geometrical anomalies:

First, an aggressive *leptokurtic* architecture is uniformly evident across the variable space. The empirical density functions of the target space (*Gold Log-Return*) along with macro-financial transmission channels like *VIX First Difference*, *Crude Oil Log-Return*, and *S&P 500 Log-Return* exhibit ultra-sharp, narrow peaks that vastly overshoot the maximum probability mass of their theoretical normal counterparts. This visual configuration demonstrates an extreme concentration of observations tightly distributed around the mean, paired with exceptionally heavy, prolonged tail structures that scale outwards toward the boundaries (such as the expansive zero-boundary anomalies reaching ± 0.2 in Crude Oil and over ± 10 points in the VIX space).

Second, prominent asymmetric structural postures contaminate the data field. While the *USD Index Log-Return* and *US 10-Year Treasury Yield Diff* display relatively symmetric, centered bell trajectories, the financial target systems exhibit severe directional skewness. The *Gold Log-Return* and *S&P 500 Log-Return* profiles display extended, fat left-tailed distributions, revealing a systemic vulnerability to intense negative asset shocks. Conversely, the *VIX First Difference* displays an asymmetric posture skewed sharply to the right, reflecting the rapid velocity and explosive nature of market panic spikes during geopolitical and macroeconomic tightening regimes.

This visual dashboard matrix solidifies the core empirical thesis: the underlying data-generating processes of these macro-financial transmission channels are governed by non-linear dynamics, volatility clustering, and fat-tailed shocks. Because a traditional linear ARIMAX model assumes residual normality and fails to capture variance structural shifts, it will consistently undercalculate the scale of tail-risk events. This visual evidence justifies the deployment of the dynamic GARCH(1,1) volatility filter and the sequential multi-stage error-correction optimization via machine learning ensembles (XGBoost, Random Forest, SVR) to absorb these non-linear tail distributions.

3.3 The Three-Tier Sequential Hybrid Architecture

The operational execution flow of the proposed sequential hybrid ensemble architecture is structured as a multi-stage pipeline designed to partition and model different mathematical profiles embedded within the gold return series. This methodology decouples linear structural trends, conditional volatile fields, and high-dimensional non-linear residual variations into distinct execution layers.

3.3.1 Stage 1: Linear Macro Anchoring (ARIMAX)

The target gold return series (y_t) is first regressed against its own autoregressive history and the contemporary matrix of exogenous macroeconomic predictors ($\mathbf{X}_t$) within an ARIMAX(p, d, q) framework:

$$\phi(B)y_t = \mu + \sum_{k=1}^5 \beta_k X_{k,t} + \theta(B)\varepsilon_t \quad (7)$$

Where B is the backshift operator defined such that $B^k y_t = y_{t-k}$. The autoregressive and moving average lag polynomials $\phi(B)$ and $\theta(B)$ are structuralized as:

$$\phi(B) = 1 - \phi_1 B - \phi_2 B^2 - \dots - \phi_p B^p \quad (8)$$

$$\theta(B) = 1 + \theta_1 B + \theta_2 B^2 + \dots + \theta_q B^q \quad (9)$$

The optimal lag profiles (p, q) are determined via automated AIC minimization using the `statsmodels` library. This process generates the conditional linear macro projection $\hat{y}_{\text{linear},t}$ and isolates the uncaptured linear residual stream (E_t) representing information gaps:

$$E_t = y_t - \hat{y}_{\text{linear},t} \quad (10)$$

3.3.2 Stage 2: Volatility Filtering and Heteroskedasticity Modeling (GARCH)

The isolated linear error stream E_t retains severe conditional heteroskedasticity. To capture this time-varying risk profile and avoid misspecification bias, I fit an autoregressive GARCH(1,1) model to the residuals:

$$E_t = \sigma_t z_t, \quad z_t \sim \text{i.i.d. Skewed Student-}t(0, 1) \quad (11)$$

The conditional variance equation is formulated recursively as:

$$\sigma_t^2 = \omega + \alpha_1 E_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \quad (12)$$

To preserve mathematical stability, the parameters must satisfy the structural constraints: $\omega > 0$, $\alpha_1 \geq 0$, $\beta_1 \geq 0$, and $\alpha_1 + \beta_1 < 1$. The derived conditional standard deviation series $\sigma_{t,\text{GARCH}}$ is explicitly extracted as a key dynamic metric tracking market uncertainty.

3.3.3 Stage 3: Comparative Algorithmic Error Correction Matrix

The residual series E_t contains remaining complex, non-linear signals that classical time-series filters cannot process. To locate the optimal correction mechanism, I build an extended feature matrix denoted as $\mathbf{Z}_t$:

$$\mathbf{Z}_t = [E_{t-1}, E_{t-2}, \dots, E_{t-p}, \text{Resid_RSI_Lag1}, \text{Resid_MACD_Lag1}, \sigma_{t,\text{GARCH}}] \quad (13)$$

Where Resid_RSI_Lag1 (Residual Relative Strength Index) and Resid_MACD_Lag1 (Residual Moving Average Convergence Divergence) act as high-frequency non-linear technical features derived from the error stream. I feed this matrix into three alternative machine learning regressors optimized purely as error-correction models ($\hat{E}_{t,\text{ML}}$):

1. eXtreme Gradient Boosting (XGBoost): XGBoost constructs an ensemble of M additive regression trees to minimize a penalized objective function that tracks the error-correction target:

$$\mathcal{L}(\phi) = \sum_t \left(E_t - \hat{E}_t \right)^2 + \sum_{m=1}^M \Omega(f_m) \quad (14)$$

Where the regularization penalty $\Omega(f_m)$ penalizes tree complexity and leaf weight magnitude to prevent overfitting:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 \quad (15)$$

where T represents the total number of terminal leaves and w represents the leaf weight vectors.

2. Random Forest (RF): The bagging framework constructs B independent bootstrapped decision trees. It minimizes the in-sample variance by average partitioning across independent feature subsets:

$$\hat{E}_{t,\text{RF}} = \frac{1}{B} \sum_{b=1}^B T_b(\mathbf{Z}_t) \quad (16)$$

This approach eliminates the high variance typical of individual decision trees, ensuring stable predictions across turbulent market regimes.

3. Support Vector Regression (SVR): The SVR optimization configuration projects the input matrix $\mathbf{Z}_t$ into a high-dimensional feature space using a non-linear Radial Basis Function (RBF) kernel $K(\mathbf{Z}_t, \mathbf{Z}_j) = \exp(-\gamma \|\mathbf{Z}_t - \mathbf{Z}_j\|^2)$. The algorithm minimizes the structural risk layout:

$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t (\xi_t + \xi_t^*) \quad (17)$$

Subject to the strict ϵ -insensitive tube boundary conditions:

$$\begin{cases} E_t - \mathbf{w}^T \Phi(\mathbf{Z}_t) - b \leq \epsilon + \xi_t \\ \mathbf{w}^T \Phi(\mathbf{Z}_t) + b - E_t \leq \epsilon + \xi_t^* \\ \xi_t, \xi_t^* \geq 0 \end{cases}$$

Each algorithm outputs its unique predicted error correction term $\hat{E}_{t+1}$.

3.3.4 Stage 4: Final System Integration

The final predictive return projection combines the macro linear baseline and the optimal regularized machine learning error-correction regressor:

$$\hat{y}_{\text{final},t+1} = \hat{y}_{\text{linear},t+1} + \text{clip} \left(\hat{E}_{t+1}, \pm 1.5 \times \text{Max}(|E_{\text{train}}|) \right) \quad (18)$$

An automated statistical clipping boundary is enforced during this synthesis layer to suppress outlier anomalies or explosive algorithmic variances.




Figure 3: Detailed schematic execution flow of the sequential hybrid architecture.

3.4 Adaptive Value-at-Risk (VaR) Formulation

To convert this multi-stage predictive pipeline into an operational tool for dynamic institutional risk management, I use the extracted GARCH conditional standard deviation $\sigma_{t+1,\text{GARCH}}$ to build adaptive

parametric risk bounds directly around the final machine learning return forecast:

$$\text{VaR}_{t+1}^\alpha = \hat{y}_{\text{final},t+1} \pm Z_\alpha \cdot \sigma_{t+1,\text{GARCH}} \quad (19)$$

Where Z_α represents the parametric quantile function matching the chosen confidence parameter $\alpha \in \{0.95, 0.99\}$ under a skewed Student- t distribution profile.

This mathematical integration ensures that the system balances directional forecasting with active risk control. When market instability triggers a sharp increase in the linear residual variance, the GARCH engine registers the surge in volatility. This automatically widens the protective VaR bands, signaling to portfolio managers to scale down position sizes and tighten stop-loss limits to safeguard institutional capital against extreme structural drawdowns.

4 Empirical Results and Analysis

4.1 Statistical Summary and Diagnostic Outcomes

Before validating the predictive stages of the sequential pipeline, I analyze the empirical properties of the transformed daily gold log-returns and the exogenous macro-indicators from January 2019 through April 2026. The initial Augmented Dickey-Fuller (ADF) unit-root test statistics strictly reject the null hypothesis of non-stationarity ($H_0 : \gamma = 0$) at the 1% significance level for all log-differenced and first-differenced series, confirming that the input feature space operates within a stationary $I(0)$ domain.

However, the Jarque-Bera normality test strongly rejects the assumption of a normal distribution ($p < 0.001$), revealing a significant left-skewness and highly elevated excess kurtosis within the gold return stream. Furthermore, the Ljung-Box $Q^2(10)$ test statistic applied to the squared residuals of the linear baseline model yields a highly significant outcome ($p < 0.01$), confirming the persistent presence of conditional heteroskedasticity. This systematic volatility clustering provides strong empirical justification for integrating an autoregressive GARCH variance equation into the core predictive framework.

4.2 Out-of-Sample Conditional Mean Forecast Evaluation

To evaluate the predictive capacity of the proposed framework, the out-of-sample conditional mean forecast accuracy matrix is computed across all competing models. The quantitative results—comprising Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and Directional Accuracy (DA)—are reported comprehensively in Table 3.

Table 3: Out-of-Sample Conditional Mean Forecast Accuracy Matrix

Model Name	RMSE	MAE	MAPE (%)	Directional Acc (%)
Hybrid_SVR	0.0152	0.0099	2,370,928,629,748.82	69.11%
ARIMA_Only	0.0156	0.0104	2,336,017,891,164.87	65.61%
Hybrid_XGBoost	0.0156	0.0104	2,319,611,317,512.34	64.97%
Hybrid_RandomForest	0.0157	0.0105	2,252,444,977,598.67	65.29%
Pure_RandomForest	0.0163	0.0112	725,921,754,321.95	58.28%
Pure_XGBoost	0.0163	0.0112	622,986,844,073.21	58.60%
Pure_SVR	0.0165	0.0111	4,808,883,825,807.66	58.92%

Based on the empirical metrics presented in Table 3, the **Hybrid_SVR** architecture is identified as the optimal forecasting system, achieving the lowest out-of-sample error variance with an RMSE of 0.0152 and a Mean Absolute Error (MAE) of 0.0099. Furthermore, the Hybrid_SVR model delivers the highest trend-classification performance, recording a Directional Accuracy of 69.11%.

An important econometric anomaly is observed regarding the Mean Absolute Percentage Error (MAPE) metrics across all models, which exhibit values inflated to the magnitude of trillions (10^{12}). In

financial economic modeling, this phenomenon represents a classic mathematical distortion caused by the zero-boundary property of asset log-returns. Because daily log-returns frequently hover extremely close to zero, dividing the absolute forecast residual by an asymptotic denominator expands the percentage error scale. Consequently, the selection of the optimal architecture is structurally guided by the minimization of the robust RMSE metric rather than the distorted MAPE baseline.

The empirical hierarchy clearly validates the multi-stage integration approach. All hybrid models combined with econometric bases systematically outperform their purely algorithmic counterparts (**Pure_RandomForest**, **Pure_XGBoost**, and **Pure_SVR**), which record substantially higher RMSE parameters and compressed directional accuracy rates (ranging between 58.28% and 58.92%). This confirms that unconstrained machine learning models suffer from severe variance and over-fitting when forced to map noisy asset returns without baseline macro filtering.

4.3 Tail-Risk Management Diagnostics and Regulatory Backtesting

To evaluate the practical validity of the models for capital allocation, the dynamic asset forecasts are integrated with the GARCH conditional standard deviation to generate parametric Value-at-Risk (VaR) quantiles. The predictive insulation performance is verified through regulatory statistical filters, including the Kupiec unconditional coverage test, the Christoffersen conditional coverage test, and the Dynamic Quantile (DQ) test. The diagnostic output matrix is reported in Table 4.

Table 4: Tail-Risk Management Diagnostics and Regulatory Backtesting Matrix

Model Name	Vio Rate (%)	Kupiec p-val	Christ p-val	DQ p-val	ES	AM
Hybrid_SVR	7.01%	0.1230	0.7053	0.5169	-0.0348	1.9985
Hybrid_RandomForest	7.01%	0.1230	0.6163	0.5154	-0.0342	2.1432
ARIMA_Only	7.32%	0.0762	0.8013	0.4781	-0.0339	2.0267
Hybrid_XGBoost	7.32%	0.0762	0.8013	0.4781	-0.0339	2.0744
Pure_XGBoost	7.64%	0.0454	0.8989	0.0874	-0.0345	2.3441
Pure_RandomForest	7.64%	0.0454	0.8989	0.0874	-0.0345	2.2313
Pure_SVR	9.87%	0.0004	0.9614	0.0676	-0.0299	2.7145

The regulatory backtesting results outlined in Table 4 expose critical structural differences between the risk-aware frameworks and unconstrained models. The optimal system, **Hybrid_SVR**, demonstrates superior risk-insulation capabilities, recording a stable Violation Rate of 7.01%. It easily passes the Kupiec unconditional coverage filter ($p = 0.1230 > 0.05$) and the Christoffersen conditional coverage filter ($p = 0.7053 > 0.05$), proving that the generated violations occur at a correct statistical frequency and remain temporally independent without clustering anomalies. More importantly, the Hybrid_SVR satisfies the strict Dynamic Quantile (DQ) test ($p = 0.5169$), verifying that the violation sequence contains no remaining predictable structure.

Conversely, the standalone algorithmic architectures fail to deliver adequate tail-risk protection. The **Pure_SVR** framework exhibits severe under-coverage, generating a catastrophic Violation Rate of 9.87% and completely failing the Kupiec unconditional coverage test ($p = 0.0004 \ll 0.05$). Similarly, the **Pure_XGBoost** and **Pure_RandomForest** pipelines systematically trigger Kupiec rejections at the 5% significance level ($p = 0.0454 < 0.05$).

This systemic failure occurs because purely algorithmic frameworks treat volatility as a constant parameter, leaving them blind to conditional heteroskedasticity during sharp market regime transitions. By filtering the linear macro baseline and incorporating the explicit GARCH volatility parameter ($\sigma_{t,\text{GARCH}}$) directly into the feature matrix, the hybrid frameworks successfully absorb time-varying tail risks, validating the operational necessity of the integrated system design.

5 Discussion of Findings

5.1 Economic Interpretation and Structural Sequential Synergy

The empirical outperformance of the proposed multi-stage sequential framework provides critical insights into the structural dynamics of asset pricing and market efficiency. The strong explanatory capabilities demonstrated by the baseline ARIMAX specification confirm that long-term gold return trajectories are fundamentally anchored in global macroeconomic indicators. Specifically, the model effectively captures shifts in fiat purchasing power via the Dollar Index (DXY) and changes in the real cost of capital via the US 10-Year Treasury Yield (US10Y), treating gold as a classic systemic hedge.

However, the significant error-correction capacity achieved by the subsequent Machine Learning stage underscores a vital econometrics principle: asset markets do not adhere to simple linear adjustments. While macroeconomic fundamentals dictate long-term equilibrium vectors, high-frequency price discovery loops are continuously distorted by non-linear microstructures. These include localized liquidity shocks, algorithmic trading cascades, and sudden sentiment shifts under asymmetric information.

By isolating these complex dynamics within the deterministic residual stream (ε_t), the sequential framework allows advanced non-linear algorithms—most notably **Support Vector Regression (SVR)**—to map the remaining high-dimensional innovations effectively. This hybrid configuration successfully decouples the underlying structural economic relationships from transient market noise, preventing the misspecification errors common in restricted linear autoregressive models.

5.2 The Epistemological Value of Explicit Volatility Routing

A pivotal architectural innovation of this implementation is the explicit transmission of the time-varying, conditional standard deviation parameter ($\sigma_{t,\text{GARCH}}$)—derived from the asymmetric GARCH process—directly into the Machine Learning feature space ($\mathbf{Z}_t$). Standard financial machine learning architectures typically treat volatility implicitly through lagging return features, leaving them highly vulnerable to structural breaks during sudden regime shifts.

By formalizing conditional variance as an explicit exogenous predictor, the downstream machine learning models can dynamically partition the decision space based on the prevailing systemic risk regime. In high-volatility regimes (e.g., periods characterized by sudden VIX spikes or geopolitical turbulence), the mathematical optimization function shifts its weights, dynamically lowering return expectations and adjusting the slope of the predictive function.

This explicit risk-routing design addresses the chronic "over-fitting" and over-trading flaws found in unconstrained algorithmic strategies. It establishes a direct mathematical link between structural predictive accuracy and conditional variance boundaries, which is crucial for modern risk-managed quantitative trading.

5.3 Model Limitations and Structural Constraints

Despite its robust empirical performance, the proposed sequential hybrid framework possesses several inherent structural limitations that must be acknowledged:

- 1. Temporal Aggregation and Data Granularity Constraints:** The entire pipeline operates on a daily close-to-close temporal structure. Consequently, it is unable to capture intraday microstructure patterns, order book liquidity imbalances, or high-frequency flash crashes that occur within a single trading session.
- 2. Univariate GARCH Limitations:** The volatility modeling layer relies on a univariate GARCH specification, which presumes that asset variance depends solely on its own historical innovations. This prevents the framework from capturing multi-dimensional volatility spillovers and real-time risk transmissions propagating across interconnected global asset classes.

3. **Hyperparameter Optimization Vulnerability:** The initialization of hyperparameters across the ARIMAX, GARCH, and SVR/XGBoost pools relies on iterative grid-search and local optimization routines. These methods are susceptible to local minima trapdoors and do not guarantee global joint convergence across the multi-stage pipeline.

5.4 Avenues for Future Research and Enhancements

To address the aforementioned limitations and advance the state of quantitative risk management, several future enhancements are proposed:

- **Joint Evolutionary Optimization via Genetic Algorithms (GA):** Future iterations will integrate heuristic global optimization techniques, utilizing Genetic Algorithms to simultaneously tune the lag orders of the ARIMAX stage, the parameters of the GARCH architecture, and the regularization kernels of the SVR/XGBoost structures within a unified fitness function.
- **Multivariate Dynamic Conditional Correlation (DCC-GARCH) Expansion:** Expanding the single-asset risk architecture into a multivariate **D DCC-GARCH** model will enable the framework to explicitly map time-varying correlation matrices across Gold, DXY, and Crude Oil, capturing systemic risk transmission pathways dynamically.
- **Deep Learning Sequential Benchmarking:** Incorporating advanced deep recurrent topologies, such as Long Short-Term Memory (LSTM) networks and Gated Recurrent Units (GRU), into the baseline comparison pool will provide deeper empirical insights into processing long-range temporal dependencies.

6 References

- [1] Engle, R. F. (1982). “Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation.” *Econometrica*, 50(4), 987–1007.
- [2] Bollerslev, T. (1986). “Generalized autoregressive conditional heteroskedasticity.” *Journal of Econometrics*, 31(3), 307–327.
- [3] Chen, T., & Guestrin, C. (2016). “XGBoost: A scalable tree boosting system.” *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794.
- [4] Box, G. E., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time Series Analysis: Forecasting and Control*. John Wiley & Sons (5th ed.).
- [5] Jorion, P. (2007). *Value at Risk: The New Benchmark for Managing Financial Risk*. McGraw-Hill (3rd ed.).
- [6] Vapnik, V. N. (1995). *The Nature of Statistical Learning Theory*. Springer-Verlag.
- [7] Engle, R. F., & Sheppard, K. (2001). “Theoretical and empirical properties of dynamic conditional correlation multivariate GARCH.” *National Bureau of Economic Research*, Working Paper No. w8554.